

---

Course : Machine Learning

Course Code : 2CS501

Machine Learning Assignment

▼ *Member 1: 19BCE043*

*Name: Yash Darji*

*Member 2: 19BCE069*

*Name: Ayush Gohil*

▼ Dataset number - 4

Adult Data Set

Dataset Description: Predict whether income exceeds \$50K/yr based on census data. Also known as "Census Income" dataset.

Link : <http://archive.ics.uci.edu/ml/datasets/Adult>

1. Download adult.data from <http://archive.ics.uci.edu/ml/machine-learning-databases/adult/>
2. save as csv file
3. read as csv
4. same steps for adult.test file
5. training dataset in dataframe df
6. test dataset in dataframe df\_test
7. importing libraries
8. printing dataset head and information
9. But dataset is not contain any column/attribute information
10. So now we are adding attributes name / column's name
11. Now dataframes are with column and headings
12. now for each attribute we will count number of categories
13. we are counting that how many are married, divorcee, widow all that count by marital status attribute..
14. And this same for all the attributes
15. Now we are comparing two attributes one by one in compare and checking which one to ignore.
16. We can also skip Gain as we see this graph
17. Loss is also not providing much information here
18. So from all this histplots we came to conclusion that we can exclude Gain, Loss, Final Weight, Country , HPR, and race.
19. We can ignore this fetures. So that we can get clean data.
20. So we are dropping some atributes with low information...

21. drop\_columns = ['Gain', 'Loss', 'Final-weight', 'Country', 'Hours-per-week', 'Race']
22. we can remove the rows that contain one or more missing values.
23. Splitting of training and testing data
24. Now we are converting >50k in 1 as output
25. and <=50k in 0 as output as we are predicting binary 1 and 0
26. Still our data is in categorical , integer form
27. So our task is to organize it in numerical form
28. So that we can fit different models to it..
29. So using Onehot encoder and pipeline process we converted our data in well organized form successfully
30. So , Finally our data is ready to fit models to it..
31. We will use following models now one by one:

1. KNN Classifier
2. Gaussian Naive Bayes
3. Bernouli Naive Bayes
4. SVC
5. SVC with rbf kernel
6. Decision Tree
7. LogisticRegression

32. The Model KNN Has Achieved 83.06 Percent Accuracy
33. The Model Gaussian Naive Bayes Has Achieved 59.75 Percent Accuracy, Which is very low..
34. The Model Bernouli Naive Bayes Has Achieved 76.17 Percent Accuracy
35. The Model SVC Naive Bayes Has Achieved 83.22 Percent Accuracy
36. The Model SVC with kereneel rbf Has Achieved 83.15 Percent Accuracy
37. The Model Decision Tree Classifier Has Achieved 82.44 Percent Accuracy
38. The Model Logistic Regression Has Achieved 82.97 Percent Accuracy
39. So at last we are comparing different models accuracy with bar plot..
40. We successfully predicted the Income of adults, which was our dataset's objective...

```
#importing different python Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os
import csv
import seaborn as sns
```

```
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn.linear_model import LogisticRegression
# ...

from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
```

```
from requests import get
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline
from sklearn.ensemble import VotingClassifier, RandomForestClassifier
```

```
from sklearn import metrics, tree, datasets, svm
from sklearn.metrics import classification_report, confusion_matrix
```



```
#download adult.data from http://archive.ics.uci.edu/ml/machine-learning-databases/adult/
#save as csv file
#read as csv
df = pd.read_csv('adult.data.csv')
```

```
#download adult.test from http://archive.ics.uci.edu/ml/machine-learning-databases/adult/
#save as csv file
#read as csv
df_test = pd.read_csv('adult.test.csv')
```

```
print(df.shape)
#shape
```

```
(17188, 15)
```

```
print(df_test.shape)
```

```
(16281, 1)
```



```
df.isna().sum()
```

```
39      0
State-gov  0
77516    0
Bachelors  0
13       0
```

```
Never-married    0
Adm-clerical     0
Not-in-family    0
White            0
Male             0
2174             0
0                0
40               0
United-States    0
<=50K           1
dtype: int64
```

df.head()

	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
0	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
1	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
2	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
3	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K
4	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States	<=50K

df\_test.head()

1x3 Cross validator														
25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United-States	<=50K.
38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United-States	<=50K.
28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	United-States	>50K.
44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	United-States	>50K.
18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0	0	30	United-States	<=50K.

Now we will give column heading to our dataframe

▼ -----

```
#column names missing

columns = ['Age', 'Working class', 'Final-weight', 'Education', 'Education-num', 'Marital-status', 'Occupation', 'Relationship', 'Race', 'Sex', 'Gain', 'Loss', 'Hours-per-week', 'Country', 'Income']
```

Training dataframe column names

```
filename='adult.data.csv'
first_row = False
with open(filename, 'r') as file:
    reader = csv.reader(file, delimiter=',')

    with open('temp.csv', 'w') as temp:
        writer = csv.writer(temp, delimiter=',')

        line_count = 0
        for row in reader:
            if line_count == 0:
                writer.writerow(columns)
                if first_row:
                    writer.writerow(row)
            else:
                writer.writerow(row)
            line_count += 1
os.remove(filename)
os.rename('temp.csv', filename)
```

### Testing dataframe column names

```
filename='adult.test.csv'
first_row = False
with open(filename, 'r') as file:
    reader = csv.reader(file, delimiter=',')

    with open('temp.csv', 'w') as temp:
        writer = csv.writer(temp, delimiter=',')

        line_count = 0
        for row in reader:
            if line_count == 0:
                writer.writerow(columns)
                if first_row:
                    writer.writerow(row)
            else:
                writer.writerow(row)
            line_count += 1
os.remove(filename)
os.rename('temp.csv', filename)
```

```
df = pd.read_csv('adult.data.csv')
df_test = pd.read_csv('adult.test.csv')
```

```
df.head()
```

	Age	Working class	Final-weight	Education	Education-num	Marital-status	Occupation	Relationship	Race	Sex	Gain	Loss	Hours-per-week	Country	Income
0	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
1	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
2	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
3	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K
4	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States	<=50K

df\_test.head()

	Age	Working class	Final-weight	Education	Education-num	Marital-status	Occupation	Relationship	Race	Sex	Gain	Loss	Hours-per-week	Country	Income
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United-States	<=50K.
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United-States	<=50K.
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	United-States	>50K.
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	United-States	>50K.
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0	0	30	United-States	<=50K.

Now dataframes are with column and headings

now for each attribute we will count number of categories

```
#Age continues values
df['Working class'].value_counts()
#Final weights also continues values
```

```
Private      22696
Self-emp-not-inc  2541
Local-gov    2093
?            1836
State-gov    1297
Self-emp-inc  1116
Federal-gov   960
Without-pay   14
Never-worked   7
Name: Working class, dtype: int64
```

```
df['Education'].value_counts()
```

```
HS-grad      10501
Some-college  7291
Bachelors    5354
```

```
Masters          1723
Assoc-voc        1382
11th             1175
Assoc-acdm       1067
10th             933
7th-8th          646
Prof-school      576
9th              514
12th             433
Doctorate        413
5th-6th          333
1st-4th          168
Preschool        51
Name: Education, dtype: int64
```

```
df['Education-num'].value_counts()
```

```
9      10501
10     7291
13     5354
14     1723
11     1382
7      1175
12     1067
6       933
4       646
15      576
5       514
8       433
16      413
3       333
2       168
1        51
Name: Education-num, dtype: int64
```

```
df['Marital-status'].value_counts()
```

```
Married-civ-spouse    14976
Never-married         10682
Divorced               4443
Separated              1025
Widowed                993
Married-spouse-absent  418
Married-AF-spouse      23
Name: Marital-status, dtype: int64
```

```
df['Occupation'].value_counts()
```

```
Prof-specialty    4140
Craft-repair      4099
Exec-managerial   4066
Adm-clerical      3769
Sales             3650
Other-service     3295
Machine-op-inspct 2002
?                1843
Transport-moving  1597
Handlers-cleaners 1370
```

```
Farming-fishing      994
Tech-support         928
Protective-serv      649
Priv-house-serv      149
Armed-Forces         9
Name: Occupation, dtype: int64
```

```
df['Relationship'].value_counts()
```

```
Husband      13193
Not-in-family  8304
Own-child     5068
Unmarried     3446
Wife          1568
Other-relative  981
Name: Relationship, dtype: int64
```

```
df['Race'].value_counts()
```

```
White      27815
Black       3124
Asian-Pac-Islander  1039
Amer-Indian-Eskimo  311
Other       271
Name: Race, dtype: int64
```

```
df['Sex'].value_counts()
```

```
Male      21789
Female    10771
Name: Sex, dtype: int64
```

```
df['Country'].value_counts()
```

```
United-States      29169
Mexico              643
?                  583
Philippines        198
Germany            137
Canada             121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba               95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic  70
Vietnam            67
Guatemala          64
Japan              62
Poland             60
Columbia           59
Taiwan             51
```




Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Hungary	13
Honduras	13
Scotland	12
Holand-Netherlands	1

Name: Country, dtype: int64

Now we are comparing two attributes one by one in compare and checking which one to ignore.


```
sns.histplot(df, x = df['Income'], hue = df['Sex'], multiple = 'dodge');
```




```
sns.histplot(df, x = df['Age'], hue = df['Income'], multiple = 'dodge', binwidth=5);
```




```
sns.histplot(df, x = df['Occupation'], hue = df['Income'], multiple = 'dodge');
```



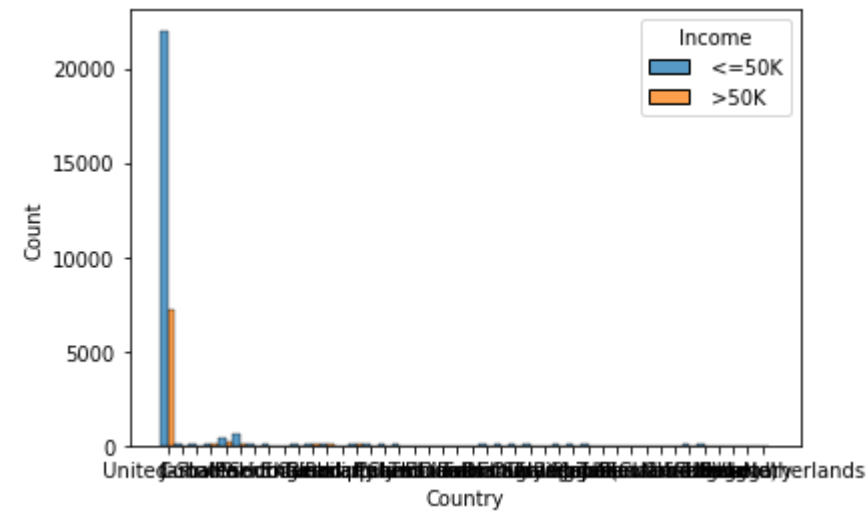
```
sns.histplot(df, x = df['Education'], hue = df['Income'], multiple = 'dodge');
```




```
sns.histplot(df, x = df['Education-num'], hue = df['Income'], multiple = 'dodge', binwidth=3);
```



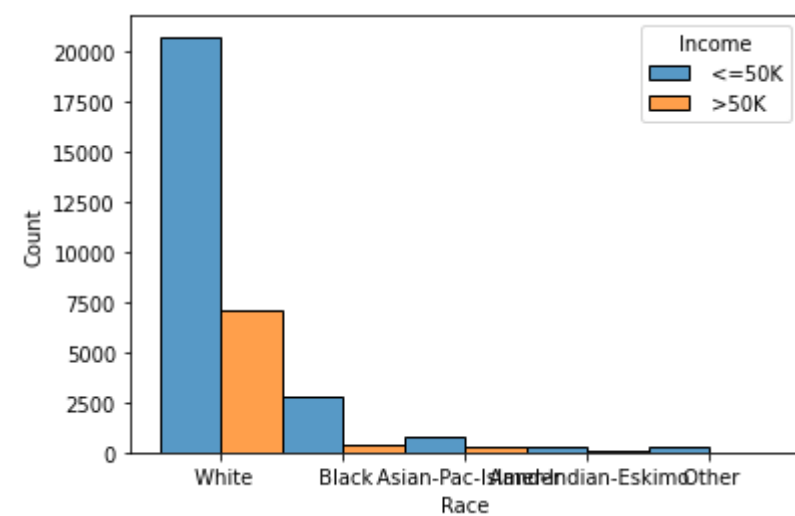
```
sns.histplot(df, x = df['Country'], hue = df['Income'], multiple = 'dodge', binwidth=3);
#we can exclude country attribute
```




```
sns.histplot(df, x = df['Hours-per-week'], hue = df['Income'], binwidth = 5, multiple = 'dodge');
```



```
sns.histplot(df, x = df['Race'], hue = df['Income'], multiple = 'dodge');
#we can skip Race
```




```
sns.histplot(df, x = df['Relationship'], hue = df['Income'], multiple = 'dodge');
```




```
sns.histplot(df, x = df['Gain'], hue = df['Income'], multiple = 'dodge');
```

#We can also skip Gain as we see this graph



```
sns.histplot(df, x = df['Loss'], hue = df['Income'], multiple = 'dodge');
```

#Loss is not providing much information here  
#We can also skip Loss as we see this graph



▼ So from all this histplots we came to conclusion that we can exclude Gain, Loss, Final Weight, Country , HPR, and race. We can ignore this fetures. So that we can get clean data.

So we are dropping some atributes with low information...

```
drop_columns = ['Gain', 'Loss', 'Final-weight', 'Country', 'Hours-per-week', 'Race']
df.drop(axis = 0, columns = drop_columns, inplace = True)
df_test.drop(columns = drop_columns, inplace = True)
```

```
df.head()
```

	Age	Working class	Education	Education-num	Marital-status	Occupation	Relationship	Sex	Income
0	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	Male	<=50K
1	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	Male	<=50K
2	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Male	<=50K
3	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Female	<=50K
4	37	Private	Masters	14	Married-civ-spouse	Exec-managerial	Wife	Female	<=50K

```
df.describe()
```

	Age	Education-num
count	32560.000000	32560.000000
mean	38.581634	10.080590
std	13.640642	2.572709
min	17.000000	1.000000
25%	28.000000	9.000000
50%	37.000000	10.000000
75%	48.000000	12.000000
max	90.000000	16.000000

```
df_test.head()
```

	Age	Working class	Education	Education-num	Marital-status	Occupation	Relationship	Sex	Income
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Male	<=50K.
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	Male	<=50K.

```
df_test.describe()
```

	Age	Education-num
count	16281.000000	16281.000000
mean	38.767459	10.072907
std	13.849187	2.567545
min	17.000000	1.000000
25%	28.000000	9.000000
50%	37.000000	10.000000
75%	48.000000	12.000000
max	90.000000	16.000000

```
print(df.shape)
print(df_test.shape)
```

```
(32560, 9)
(16281, 9)
```

```
#we can remove the rows that contain one or more missing values.
df = df.dropna()
df_test = df_test.dropna()
```

```
print(df.shape)
print(df_test.shape)
#0 rows with missing values here
```

```
(32560, 9)
(16281, 9)
```

### Splitting of training and testing data

```
#Splitting of training and testing data
train_data = df.drop(columns = ['Income'])
y_train = df['Income']

test_data = df_test.drop(columns = ['Income'])
y_test = df_test['Income']
```

```
test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16281 entries, 0 to 16280
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              16281 non-null  int64
1   Working class    16281 non-null  object
2   Education        16281 non-null  object
3   Education-num    16281 non-null  int64
4   Marital-status   16281 non-null  object
5   Occupation       16281 non-null  object
6   Relationship     16281 non-null  object
7   Sex              16281 non-null  object
dtypes: int64(2), object(6)
memory usage: 1017.7+ KB
```

```
train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32560 entries, 0 to 32559
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              32560 non-null  int64
1   Working class    32560 non-null  object
2   Education        32560 non-null  object
3   Education-num    32560 non-null  int64
4   Marital-status   32560 non-null  object
5   Occupation       32560 non-null  object
6   Relationship     32560 non-null  object
7   Sex              32560 non-null  object
dtypes: int64(2), object(6)
memory usage: 2.2+ MB
```

Now we will convert income column in 1 or 0

```
y_train_n = np.array(y_train)
y_test = np.array(y_test)

y_train_n[y_train_n == ' >50K'] = 1
y_test[y_test == ' >50K.'] = 1

y_train_n[y_train_n == ' <=50K'] = 0
y_test[y_test == ' <=50K.'] = 0

y_train_n = np.array(y_train_n, dtype = np.uint8)
y_test = np.array(y_test, dtype = np.uint8)

y_test
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:5: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise compariso
"""
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:8: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise compariso
```

```
array([0, 0, 1, ..., 0, 0, 1], dtype=uint8)
```

```
np.count_nonzero(y_train)
```

```
32560
```

Now our data is in categorical and integer form..

So our task is to transform it into organized way so that we can fit different models to it.

Some content of the next portion of code is taken from another resource for transforming dataset in organized way..

```
class MostFrequentImputer(BaseEstimator, TransformerMixin):
    def fit(self, X, y=None):
        self.mf = {}
        for col in X.columns:
            self.mf[col] = X[col].value_counts().index[0]
            self.data = X
        return self
    def transform(self, X, y=None):
        for col in X.columns:
            X[col] = X[col].replace([' ?'], self.mf[col])
        return X
```

```
Categorical_Pipeline = Pipeline(
    [
        ('Replace "?" values', MostFrequentImputer()),
        ('Encode Values', OneHotEncoder(sparse = False)),
    ]
)
```

```
Full_Pipeline = ColumnTransformer(
    [
        ('Numerical Data', StandardScaler(), ['Age', 'Education-num']),
        ('Categorical Data', Categorical_Pipeline, ['Working class', 'Education', 'Marital-status', 'Occupation',
                                                    'Relationship', 'Sex']),
    ]
)
```

```
train_data_processed = Full_Pipeline.fit_transform(train_data)
X_test = Full_Pipeline.fit_transform(test_data)
```

```
y_test
```

```
array([0, 0, 1, ..., 0, 0, 1], dtype=uint8)
```



```
X_train, X_valid, y_train, y_valid = train_test_split(train_data_processed, y_train_n, test_size = 0.4)
```

```
len(X_train), len(X_valid)
```

```
(19536, 13024)
```

So , Finally our data is ready to fit models to it..

```
from sklearn import metrics,tree,datasets,svm
from sklearn.metrics import classification_report,confusion_matrix
```

We will use following models now one by one:

1. *KNN Classifier*

2. *Gaussian Naive Bayes*

3. *Bernouli Naive Bayes*

4. *SVC*

5. *SVC with rbf kernel*

6. *Decision Tree*

7. *LogisticRegression*

## ▼ 1. KNN Classifier

I have used Grid search for getting best result among different values of hyperparameters

```
# Parameter Space of the KNN Classifier
para = [{'n_neighbors': [5, 7, 9, 10,11, 15,17 ,21,23,25]}]
knn_ = KNeighborsClassifier()
rknn = GridSearchCV(knn_, para)
rknn.fit(X_train, y_train)

GridSearchCV(cv=None, error_score=nan,
             estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                           metric='minkowski',
                                           metric_params=None, n_jobs=None,
                                           n_neighbors=5, p=2,
                                           weights='uniform'),
             iid='deprecated', n_jobs=None,
             param_grid=[{'n_neighbors': [5, 7, 9, 10, 11, 15, 17, 21, 23,
                                           25]}],
```

```
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=None, verbose=0)
```

```
rknn.best_estimator_    # Best Hyperparameters
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=None, n_neighbors=25, p=2,
                     weights='uniform')
```

```
model_knn = KNeighborsClassifier(n_neighbors=25)    # Using the best parameters
model_knn.fit(X_train, y_train)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=None, n_neighbors=25, p=2,
                     weights='uniform')
```

```
print(model_knn.score(X_train, y_train))
```

```
print(model_knn.score(X_valid, y_valid))
```

```
knn_accuracy = 100*model_knn.score(X_valid, y_valid)
```

```
0.8429054054054054
0.8306203931203932
```

```
print('The Model KNN Has Achieved %.2f Percent Accuracy'%(knn_accuracy))
```

```
The Model KNN Has Achieved 83.06 Percent Accuracy
```

*The Model KNN Has Achieved 83.06 Percent Accuracy*

## ▼ 2. Gaussian Naive Bayes

```
model_nb = GaussianNB()
model_nb.fit(X_train, y_train)
```

```
GaussianNB(priors=None, var_smoothing=1e-09)
```

```
Gnb_accuracy=100*model_nb.score(X_valid, y_valid)
```

```
print('The Model Gaussian Naive Bayes Has Achieved %.2f Percent Accuracy'%(Gnb_accuracy))
```

```
The Model Gaussian Naive Bayes Has Achieved 59.75 Percent Accuracy
```

*The Model Gaussian Naive Bayes Has Achieved 59.75 Percent Accuracy, Which is very low..*

### ▼ 3. Bernouli Naive Bayes

```
from sklearn.naive_bayes import BernoulliNB
```

```
model_nb = BernoulliNB()  
model_nb.fit(X_train, y_train)
```

```
BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)
```

```
model_nb.score(X_valid, y_valid)
```

```
0.7616707616707616
```

```
bnb_accuracy=100*model_nb.score(X_valid, y_valid)
```

```
print('The Model Bernouli Naive Bayes Has Achieved %.2f Percent Accuracy'%(bnb_accuracy))
```

```
The Model Bernouli Naive Bayes Has Achieved 76.17 Percent Accuracy
```

*The Model Bernouli Naive Bayes Has Achieved 76.17 Percent Accuracy*

### ▼ 4. SVC

```
model=SVC()  
model.fit(X_train,y_train)  
pred=model.predict(X_valid)
```

```
print('The Model Has Achieved %.2f Percent Accuracy'%(100*metrics.accuracy_score(y_valid,pred)))  
SVM_accuracy = (100*metrics.accuracy_score(y_valid,pred))  
print(SVM_accuracy)
```

```
The Model Has Achieved 83.22 Percent Accuracy  
83.21560196560198
```

*The Model SVC Naive Bayes Has Achieved 83.22 Percent Accuracy*

### ▼ 5. SVC [kernel RBF]

```
model_sc = SVC(kernel='rbf', gamma = 0.1, C = 1) # Radial Basis Function for non-linear decision boundary
model_sc.fit(X_train, y_train)
```

```
SVC(C=1, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.1, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
model_sc.score(X_valid, y_valid)
```

```
0.831541769041769
```

```
svc_rbf_accuracy=100*model_sc.score(X_valid, y_valid)
```

```
print('The Model SVC with kereneel rbf Has Achieved %.2f Percent Accuracy'%(svc_rbf_accuracy))
```

```
The Model SVC with kereneel rbf Has Achieved 83.15 Percent Accuracy
```

*The Model SVC with kereneel rbf Has Achieved 83.15 Percent Accuracy*

## ▼ \_\_ 6. Decision Tree Classifier\_\_

```
# Parameter Space of the Decision Tree Classifier
param_dt = [
    {'max_depth': [2, 4, 7, 8, 10, 12, 15, 20, 30, 50, 100] },
    {'max_leaf_nodes': [4, 5, 7, 10, 20, 25, 30, 50, 70, 100, 150]},
]

model_dt_temp = DecisionTreeClassifier()
rsdt = GridSearchCV(model_dt_temp, param_dt, cv = 5, n_jobs = -1, return_train_score = True)
rsdt.fit(X_train, y_train)
rsdt.best_estimator_ # Best Hyperparameters
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                        max_depth=None, max_features=None, max_leaf_nodes=50,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')
```

```
model_dt = DecisionTreeClassifier(criterion='gini', max_leaf_nodes = 100) # Using the best hyperparameters
model_dt.fit(X_train, y_train)
model_dt.score(X_valid, y_valid)
```

```
0.8244011056511057
```

```
dt_accuracy=100*model_dt.score(X_valid, y_valid)
```

```
print('The Model Decision Tree Classifier Has Achieved %.2f Percent Accuracy'%(dt_accuracy))
```

The Model Decision Tree Classifier Has Achieved 82.44 Percent Accuracy

*The Model Decision Tree Classifier Has Achieved 82.44 Percent Accuracy*

## ▼ 7. Logistic Regression

```
model_lr = LogisticRegression(n_jobs = -1)
model_lr.fit(X_train, y_train)
model_lr.score(X_valid, y_valid)
```

0.8296990171990172

```
lr_accuracy=100*model_lr.score(X_valid, y_valid)
```

```
print('The Model Logistic Regression Has Achieved %.2f Percent Accuracy'%(lr_accuracy))
```

The Model Logistic Regression Has Achieved 82.97 Percent Accuracy

*The Model Logistic Regression Has Achieved 82.97 Percent Accuracy*

## ▼ Comparision of different Model's accuracy

```
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(10,6),dpi=80)
List=[knn_accuracy,Gnb_accuracy,bnb_accuracy,SVM_accuracy,dt_accuracy,lr_accuracy]
l = ['KNN', 'Gaussian', 'BERNOULI','SVC', 'Decision Tree','LogisticRegression']

c=['b','black','orange','lightpink','skyblue','maroon']
plt.bar(l,List, width = 0.6,color=c)
plt.xlabel('Machine learning Models')
plt.ylabel('Accuracy')
plt.show()
```

